

GAURAV JOSHI

AI Product Manager | Applied LLM & RAG Products | Enterprise & Regulated Environments | Zero-to-One Discovery

+91-9823466677 | kavee.gauravjoshi@gmail.com | [linkedin.com/in/gsjoshi](https://www.linkedin.com/in/gsjoshi) | gsj24983.github.io | Pune, India

PROFESSIONAL SUMMARY

Product leader, 18+ years in enterprise and regulated environments, now focused on applied AI and one of the few at this level who has shipped a multimodal RAG application solo, end to end, rather than only sponsoring one. Builds AI where the constraint is trust rather than novelty: human-in-the-loop review with audit trails, compliance risk detection grounded in regulatory source documents, and an original AI Trustworthiness Matrix built on the NIST AI RMF. Completing MIT xPRO's Designing and Building AI Products and Services, with two Exemplary Assignment awards.

AI PRODUCT WORK

- Tata Safari AI Assistant - live multimodal RAG application, built solo end to end (2026): ingestion pipeline over official service and infotainment manuals, semantic chunking with overlapping segments, semantic search over a vector store generated from the manuals, Google Gemini for multimodal reasoning across text and dashboard-warning-light images, deployed on Streamlit Cloud with session rate limiting and access control. Used as the first test case for the AI Trustworthiness Matrix. safari-ai-assistant-gsj24983.streamlit.app | github.com/GSJ24983/safari-ai-assistant
- AI Trustworthiness Matrix (personal framework): rates 8 AI-trust dimensions built on the NIST AI RMF; each gets a rating, a confidence level and the evidence actually inspected, and the action comes from the combination, not the score - a low score you are sure of and a high score you are guessing about need very different responses. Turns 'the model is probably fine' into a prioritised fix-list.
- Policy transparency and supplier content compliance (JPMC, PoC delivered): analysed detractor voice-of-customer data, supplier content and regulatory verbiage to surface 3 causal loops driving compliance risk and NPS damage. Converted them into two validated solutions - a readability redesign of lodging rules, policy and cancellation content, and automated scanning of static and API supplier content against regulatory source documents to flag risky terms and trigger action (content suppression, proactive supplier outreach). Both went to the owning teams; the scanning concept is on the AI team's backlog.
- Goal-Based Travel Package prototype: took a scoring framework out of documents and into working software, coded solo at an internal hackathon using Copilot - JavaScript scoring engine, externalised configuration, and an HTML output layer rendering each package with a visual justification of its score.
- Puppeteer (solo build, public on GitHub): branching-scenario learning engine behind Swanubhava™, my product management behavioural model - learners write free-text answers graded against a rubric, with coaching on missing criteria and competency-tagged session summaries; 97 automated tests, single-file offline export. github.com/GSJ24983/puppeteer-sbl
- LLM-assisted resume review (JPMC): designed key-term logic, synonym mapping and conflict rules, with human-in-the-loop review and a full audit trail - ~75% less review time per resume.

PRODUCT CREDIBILITY BEHIND THE AI WORK

- Five zero-to-one starts across five domains – travel e-commerce, loyalty, banking procurement, media & entertainment and enterprise software. Each with little or no data, team or precedent, including a CEO's 8-page concept taken to commercial launch in 12 months as a two-person team.
- Regulatory and governance depth that AI products need: 15+ product teams and 40 stakeholders taken to 100% dark-pattern compliance with no reporting line.
- Scale credibility: a travel and loyalty platform processing ~\$25M monthly GTV, 27-person cross-functional team, 100% client adoption on a cloud migration with zero downtime.

PROFESSIONAL EXPERIENCE

Vice President, Product Management - Travel & Cross-Product Strategy | JPMorgan Chase & Co. | Feb 2020 - Sep 2026

- Owned AI product initiatives across the travel and loyalty portfolio - use case identification, feasibility validation, evaluation criteria, human-in-the-loop design and executive business cases.
- Sole owner of cross-product greenfield initiatives with no data, benchmarks or team: scope definition, primary and secondary research, GTM strategy, feature mapping, mock screens, business case and cost-benefit analysis through to PoC decisions.
- Led product for a partner travel and loyalty platform processing ~\$25M monthly GTV across flights, hotel, car and activities, with a 27-person cross-functional team spanning product, engineering, data and UX.
- Managed 4 product managers through the cloud migration programme, each owning an independent partner workstream, with dotted-line decision authority over their engineering and architecture teams; held consolidated delivery accountability.
- Delivered 100% client adoption of the cloud loyalty redemption platform - 3 launches for 5 enterprise clients in 12 months, zero critical defects, zero downtime - enabling the firm's data-centre exit. Shipped a new UI and zero-to-one features along the way: guest flow and new forms of payment.
- Inherited a stalled supplier-error workstream (FIFT) and drove it to outcome: support alerts down 50% and booking failures down 30% across 9 suppliers, with 38 standardised error categories. Set the sequencing basis (alert volume against booking impact) and

held the resolution cadence with API and connector teams; ground-level taxonomy analysis was owned by a product analyst on the workstream.

- Led enterprise dark-pattern regulatory compliance (FTC, CA, TX, MA; equivalent to India CPA 2019 and DPDP Act 2023) across 15+ product teams and 40 stakeholders with no direct authority - built the governance model, translated regulation into product-level decisions, 100% compliant within mandated timelines. Directly relevant to AI governance, auditability and model-risk contexts.

Solution Architect / Senior Business Analyst | Cybage Software Pvt. Ltd. | May 2015 - Feb 2020

- Owned product definition for foundational platform components. Role-based access control and a configurable rules engine.
- Managed 3 senior business analysts on independent module workstreams while holding consolidated roadmap accountability; mentored 5 junior analysts through BABOK-anchored training content built in-house.
- Introduced a decision-table method for modelling complex rules and edge cases that other product teams adopted.
- Authored the white paper Smart Content Scheduling using the Analytic Hierarchy Process (AHP).

Enterprise Transformation Consultant / Product Owner | Nihilent Technologies & Medscheme | Feb 2008 - May 2015

Led 4-5 consultants and business analysts across assignments, delivered in India, South Africa and the UK.

- Customer loyalty evaluation product: took a CEO's 8-page concept to commercial launch in 12 months as a two-person team, with full go-to-market, and built an analysis framework that became a standard product feature shared with clients.
- Standard Bank South Africa: business case, cost-benefit analysis and a live 12-person procure-to-pay contact centre PoC delivered in 3 months across 5 branches, as a contractor with no formal authority. Rolled out nationally; CIPS award winner.
- Banking Association South Africa: strategy map and balanced scorecard implementation in 6 months using Nihilent's LAMAT strategy-execution platform. Metric definition and measurement design at enterprise scale. Taught as a case study at the University of the Witwatersrand.
- SuperSport and M-Net: 7 consecutive enterprise releases over 2.5 years with 2 analysts reporting in.
- Medscheme (Proxy Product Owner): backlog ownership for a healthcare core system upgrade, with 3 business analysts on independent modules.

CORE SKILLS

AI & Applied ML: AI Product Management | Generative AI | Large Language Models | Retrieval-Augmented Generation (RAG) | Multimodal RAG | Vector Search | Prompt and Evaluation Logic Design | Human-in-the-Loop Design | AI Use Case Identification | AI Trust & Risk Assessment (NIST AI RMF) | AI Governance, Auditability & Compliance

Product Leadership: Product Strategy | Zero-to-One Discovery | Product Roadmap | MVP Definition | Go-to-Market Strategy | OKRs & KPIs | Prioritisation Frameworks | Lifecycle Management

Discovery & Research: Primary and Secondary Research | Persona Development | Affinity & Multi-Criteria Scoring | Competitive Analysis | Business Case | Cost-Benefit Analysis | Success Metrics & Evaluation Design

Delivery & Leadership: Cross-Functional Leadership | Influence Without Authority | People and Workstream Management | Stakeholder Management | Agile & SAFe | Mentoring

Domain & Tools: Travel E-Commerce | Loyalty & Redemption | Banking Procurement | Enterprise Platforms | Regulatory Compliance | JavaScript, Streamlit, GitHub Copilot | Jira | Confluence | Grafana | Kibana | SAS

CERTIFICATIONS & ONGOING TRAINING

- Designing and Building AI Products and Services - Massachusetts Institute of Technology (MIT xPRO). In progress, 2026.
- SAFe Product Owner / Product Manager (POPM) - Scaled Agile | Certified Scrum Product Owner (CSPO) - Scrum Alliance | Certified Business Analysis Professional (CBAP) - IIBA

EDUCATION

P.G.D.B.M. Marketing | Indira School of Business Studies, Pune | 2006-08

B.E. Instrumentation Engineering | D. Y. Patil College of Engineering, Pune | 2001-05

AWARDS, PUBLICATIONS & RECOGNITION

- Winner, CIPS Southern Africa Best Process Improvement Initiative 2010 (Standard Bank South Africa). Outstanding Alumni Award. Nihilent Best Team Award.
- Contributor, IIBA Product Ownership Analysis Guide. Author, white paper: Smart Content Scheduling using the Analytic Hierarchy Process (AHP). Speaker, IPA Summit ('The Sherlock Way'). Featured case study, University of the Witwatersrand (WITS). 28 long-form articles published on LinkedIn.
- Two Exemplary Assignment awards during Massachusetts Institute of Technology (MIT xPRO) certification: AI-enabled organisational strategy, and a responsible-AI feasibility assessment of a generative solution.